

Terraform



Terraform

- What is Terraform and how it works (History and advantages)?
- Terraform installation
- Terraform alternatives
- Terraform - syntax.
- Terraform Commands (init, plan, apply, state list, format, validate, destroy, -auto-approve, -targets).
- Terraform - Practical.
- Terraform .tfstate file
- Creating AWS Resources (IAM, S3 Bucket, VPC)
- Variables in Terraform.



What is Terraform and how it works (History and advantages)?

What is Terraform and how it works (History and advantages)?

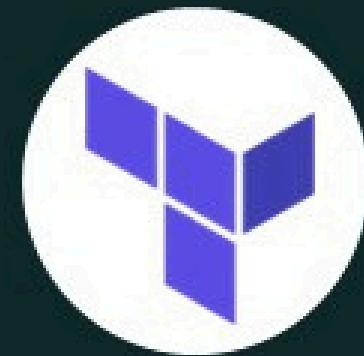


- Terraform is an Infrastructure as a Code (IAAC) tool developed by Mitchel Hashimoto in 2014.
- It was developed in “Go” language. We use HCL (HashiCorp Configuration Language) to create the infra.
- Terraform Simplified multi-cloud deployments with a single tool.



Why Terraform (advantages)?

Why Terraform (advantages)?



- Terraform is a tool used to make infrastructure automation, It is free (but, not open source), easy to understand and platform independent.
- It has many advantages..
 - ❖ Reusability
 - ❖ Time saving
 - ❖ Automation
 - ❖ Avoiding mistakes
 - ❖ Multi Cloud Support



Terraform installation

Terraform installation



Windows

- Go to the Terraform website:
<https://developer.hashicorp.com/terraform/downloads>
- Download the Windows zip for your system (usually 64-bit).
- Unzip the contents into a folder (e.g., C:\terraform)
- Add that folder to your system PATH:
- Search “Environment Variables” > Open “Edit the system environment variables”
- Click Environment Variables
- Under System variables > find “Path” > click Edit
- Click New > paste C:\terraform
- Open command prompt → terraform -v (Response: Terraform <<version>>)

Amazon Linux

- Go to the Terraform website:
<https://developer.hashicorp.com/terraform/downloads>
- <https://developer.hashicorp.com/terraform/install>
Go to Linux → package manager.
- Copy the package according to your Linux distribution.
- Go to your instance, paste and run the commands.



Terraform - syntax.



Terraform - syntax.

Syntax:

```
provider "<<provider>>" {
  region = "<<region id>>"
}

resource "<<resource type>>" "<<Resource Name (Logical Name)>>"
{
  <<Arguments (Configuration Settings)>>
}
```

Example:

```
provider "aws" {
  region = "ap-south-2"
}

resource "aws_instance" "ccitinst"
{
  ami = "ami-0a19d948cf5ce40b2"
  instance_type = "t3.micro"
}
```



Terraform – Resource components.

In Terraform, a **resource** is a fundamental building block that manages AWS infrastructure components like S3 buckets, EC2 instances, VPCs, and more.

Key Components of a Terraform Resource

1. Resource Type

The type of AWS resource being created.

Example: "aws_s3_bucket", "aws_instance", "aws_vpc".

2. Resource Name (Logical name)

A unique identifier for referencing the resource within Terraform.

Example: "my_bucket"

3. Arguments

Define properties and settings for the resource.



Terraform alternatives

Terraform alternatives





Terraform Commands.



Terraform Commands.

1. `terraform init`
2. `terraform plan`
3. `terraform apply -auto-approve`
4. `terraform state list`
5. `terraform fmt`
6. `terraform validate`
7. `terraform destroy -auto-approve`
8. `terraform destroy --auto-approve -target="aws_instance.one[1]"`



Terraform Commands.

Initialize (terraform init):

Downloads provider plugins (e.g., AWS, Azure, GCP).

Sets up the backend for storing Terraform state. It will store information of plugins in .terraform folder

Plan (terraform plan):

Shows what Terraform will create, update, or delete before applying changes.

Helps avoid accidental modifications.

Apply (terraform apply):

Executes the planned changes and provisions the infrastructure.

-auto-approve

This option allows Terraform to apply changes without requiring manual confirmation.



Terraform Commands.

terraform state list

- The “terraform state list” command displays all managed resources in the Terraform state file.
- It helps you identify the resources Terraform is tracking.

Destroy (terraform destroy):

Destroys all resources defined in the configuration.

-auto-approve

This option allows Terraform to destroy the resources without requiring manual confirmation.

terraform Validate (Validate Syntax & Configuration)

- Checks if the Terraform configuration is syntactically correct and valid.
- Detects errors like missing arguments, incorrect types, or invalid providers.



Terraform Commands.

Terraform Format – Terraform fmt

- Automatically formats Terraform configuration files (.tf and .tfvars) to follow best practices.
- Ensures consistent indentation, spacing, and style. It will be applied to all files in current folder.

--recursive

- Recursively formats Terraform configuration files in the current directory and all sub-directories.

```
provider "aws" {  
region = "us-east-1"  
}  
  
resource "aws_instance" "one" {  
count = 5  
ami = "ami-03eb6185d756497f8"  
instance_type = "t2.micro"  
}
```

```
provider "aws" {  
region = "us-east-1"  
}  
  
resource "aws_instance" "one" {  
count          = 5  
ami           = "ami-03eb6185d756497f8"  
instance_type = "t2.micro"  
}
```



Terraform Commands.

terraform Validate (Validate Syntax & Configuration)

- Checks if the Terraform configuration is syntactically correct and valid.
- Detects errors like missing arguments, incorrect types, or invalid providers.

terraform target

The `-target` flag allows you to apply or destroy specific resources instead of the entire infrastructure. This is useful when you want to update or destroy only one resource without affecting others.

Apply a specific resource:

```
terraform apply -target=aws_instance.example
```

- This updates only `aws_instance.example` without affecting other resources.

```
terraform destroy -target=aws_instance.example
```

- This deletes only the `aws_s3_bucket.my_bucket` resource.

When to Use `-target`?

- ✓ When testing changes for a single resource.
- ✓ When debugging issues with a specific resource.
- ✗ Not recommended for regular use, as it may lead to dependency issues.



Terraform - Practical.



Terraform “.tfstate” file



Terraform ".tfstate" file

- The Terraform state file (`terraform.tfstate`) is a JSON file that stores information about the infrastructure Terraform manages.

◆ Purpose of .tfstate

1. Tracks the current state of resources deployed using Terraform.
2. Maps Terraform configurations to real-world resources in AWS, Azure, etc.
3. Speeds up Terraform operations by avoiding unnecessary changes.

◆ Key Contents of .tfstate

1. Resource Details: Stores attributes like instance ID, AMI, security groups, etc.
2. Dependencies: Tracks relationships between resources.
3. Metadata: Contains information about the Terraform version used.



Terraform ".tfstate" file

- The Terraform state file (`terraform.tfstate`) is a JSON file that stores information about the infrastructure Terraform manages.

◆ Purpose of .tfstate

1. Tracks the current state of resources deployed using Terraform.
2. Maps Terraform configurations to real-world resources in AWS, Azure, etc.
3. Speeds up Terraform operations by avoiding unnecessary changes.

◆ Key Contents of .tfstate

1. Resource Details: Stores attributes like instance ID, AMI, security groups, etc.
2. Dependencies: Tracks relationships between resources.
3. Metadata: Contains information about the Terraform version used.



Terraform ".tfstate" file

◆ Important Notes

1. **Do NOT manually edit .tfstate unless necessary.**
2. **Should be stored securely** (contains sensitive data).
3. **For teams, use remote state storage** (e.g., S3, Terraform Cloud) to avoid conflicts.

```
json
{
  "resources": [
    {
      "type": "aws_instance",
      "name": "ccitinst",
      "instances": [
        {
          "attributes": {
            "ami": "ami-0a19d948cf5ce40b2",
            "instance_type": "t3.micro",
            "id": "i-1234567890abcdef"
          }
        }
      ]
    }
  ]
}
```



Terraform Comments



Terraform Comments

Terraform supports comments to improve readability and document your Infrastructure as Code (IaC).

There are two types of comments in Terraform:

1. Single-Line Comments (# or //)
2. Multi-Line Comments /* */

These comments are ignored by Terraform and do not affect execution.



Terraform Comments

Terraform supports comments to improve readability and document your Infrastructure as Code (IaC).

1. Single-Line Comments (# or //)

Used for brief comments on a single line.

```
# This is a single-line comment

// Another way to write a single-line comment

resource "aws_instance" "example" {
    ami      = "ami-12345678" # AMI ID for EC2 instance
    instance_type = "t2.micro"
}
```



Terraform Comments

Terraform supports comments to improve readability and document your Infrastructure as Code (IaC).

2. Multi-Line Comments /* ... */

Used when you need to comment out multiple lines of code.

```
/* This is a multi-line comment. It can span multiple
lines and is useful for large descriptions.

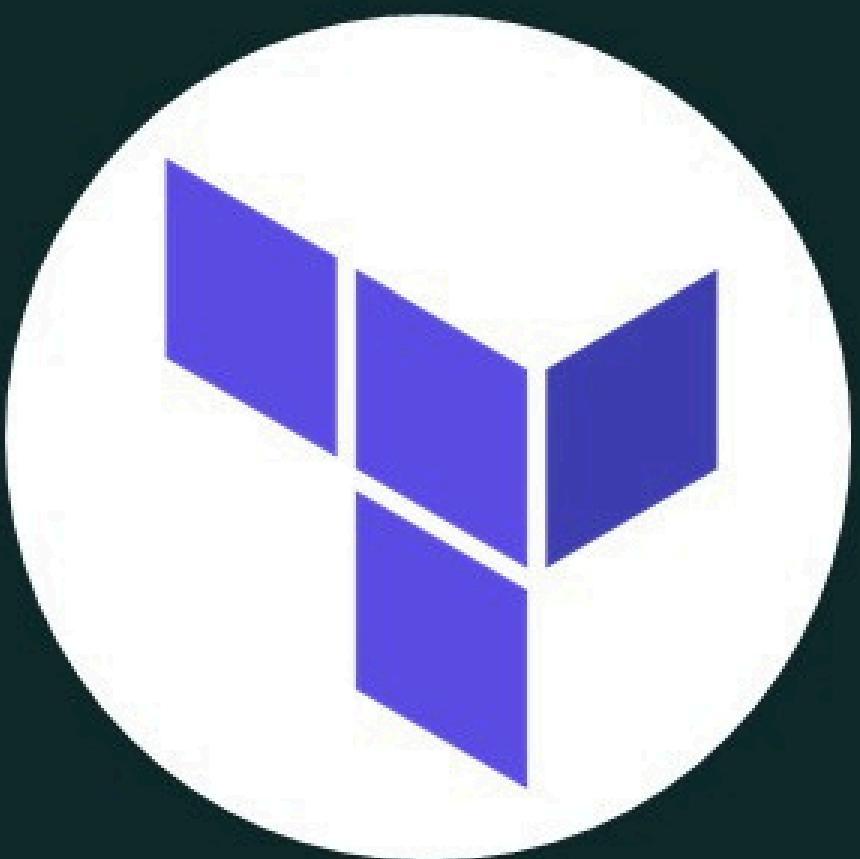
resource "aws_s3_bucket" "example" {
  bucket = "my-terraform-bucket"
}
```



Variables in Terraform.

Terraform

- Variables in Terraform.
- Terraform “.tfvars” files
- Terraform Output block
- Terraform Taint and replace
- Terraform Import
- Creating AWS Resources (EC2, EBS and RDS)





Variables in Terraform.

Terraform variables let you **customize** your infrastructure without hardcoding values. They make your code **reusable and flexible**.

Types of Variables

String → Text values

```
variable "instance_type"  
{  
  type    = string  
  default = "t2.micro"  
}
```

Number → Numeric values

```
variable "instance_count"  
{  
  type    = number  
  default = 2  
}
```

Boolean → True/False values

```
variable  
"enable_monitoring" {  
  type    = bool  
  default = true  
}
```



Variables in Terraform.

Terraform variables let you **customize** your infrastructure without hardcoding values. They make your code **reusable and flexible**.

Types of Variables

List → Ordered collection of values

```
variable "instance_types" {  
    type = list(string)  
    default = ["t2.micro", "t3.micro"]  
}
```

Map → Key-value pairs

```
variable "ami_ids" {  
    type = map(string)  
    default = {  
        "us-east-1" = "ami-12345"  
        "ap-south-1" = "ami-67890"  
    }  
}
```



Terraform “.tfvars” files



Using .tfvars Files

- Terraform allows you to use multiple .tfvars files to manage different environments like dev, staging, and production without changing the main configuration files.

Command:

```
terraform apply -var-file="prod.tfvars"
```

When to Use .tfvars files?

- ✓ Helps manage multiple environments easily.
- ✓ Prevents accidental changes in production.
- ✓ Improves Terraform configuration modularity.



Terraform Output block



Terraform Output block

- In Terraform, outputs allow you to display or pass values from your infrastructure deployment.
- Outputs are useful for retrieving details like VPC IDs, instance public IPs, S3 bucket names, and more after applying your Terraform configuration.

How to Define Outputs

```
output "vpc_id" {  
    description = "The ID of the created VPC"  
    value = aws_vpc.my_vpc.id  
}  
output "s3_bucket_name" {  
    description = "The name of the created S3  
    bucket" value = aws_s3_bucket.my_bucket.id  
}
```



Terraform Taint and replace



Terraform Taint

- It allows for you to manually mark a resource for recreation.
- in real time sometimes resources fails to create so to recreate them we use taint
- in new version we use -replace option
- TO TAINT: `terraform taint aws_instance.one[0]`
- TO UNTAINT: `terraform untaint aws_instance.one[0]`
 - `terraform state list`
 - `terraform taint aws_instance.one[0]`
 - `terraform apply --auto-approve`
- TERRAFORM REPLACE: `terraform apply --auto-approve -replace="aws_instance.one[0]"`



Terraform Import



Terraform Import

- When we create resource manually terraform wont track that resource.
- Import command can used to import the resource which is created manually.
- can only import one resource at a time (FROM CLI).
- It can import both code to config file and state file.
- FOR STATE FILE: `terraform import aws_instance.one` (not preferable).
- FOR STATEFILE & CODE:
 - `terraform plan -generate-config-out=manual.tf`
 - `terraform apply -auto-approve`

Terraform

- Terraform Debugging
- Alias & providers
- Terraform Parallelism.
- Terraform Dependency.





Terraform Debugging



Terraform Debugging

- Debugging helps understand issues (like misconfigurations or dependency errors can occur) and fix them efficiently.
- You can set TF_LOG to one of the log levels TRACE, DEBUG, INFO, WARN or ERROR to change the verbosity of the logs, with TRACE being the most verbose.





Terraform Debugging

To set LOGs:

(Linux/macOS) `export TF_LOG=TRACE` / (Windows) `$env:TF_LOG="DEBUG"`

(Linux/macOS) `export TF_LOG_PATH="logs.txt"` / (Windows) `$env:TF_LOG_PATH="log.txt"`

`terraform apply`

export TF_LOG=TRACE	export TF_LOG=DEBUG	export TF_LOG=INFO	export TF_LOG=WARN	export TF_LOG=ERROR
TRACE	DEBUG	INFO	WARN	ERROR
TRACE	DEBUG	INFO	WARN	ERROR
DEBUG	INFO	WARN	WARN	ERROR
INFO	WARN	ERROR	WARN	ERROR
WARN	ERROR			
ERROR				

To Unset LOGs:

(Linux/macOS) `unset TF_LOG` / (Windows) `Remove-Item Env:TF_LOG`

(Linux/macOS) `unset TF_LOG_PATH` / (Windows) `Remove-Item Env:TF_LOG_PATH`



Terraform

Alias and Providers



Alias and Providers

- In Terraform, providers are responsible for managing and interacting with external services (like AWS, Azure, GCP, etc.).
- Aliases allow you to define multiple configurations of the same provider within a single Terraform configuration.

What is a Provider?

- A provider is a plugin that enables Terraform to interact with an external service.

```
provider "aws"  
{  
    region = "ap-south-2"  
}
```

- Here, Terraform downloads the AWS provider and sets the region to ap-south-2.



Why Use an Alias?

- The alias argument allows multiple configurations of the same provider. This is useful when:
 1. You can deploy resources in multiple AWS regions.
 2. You can deploy in multiple AWS accounts.
 3. You can deploy to different service providers.



Terraform

1. You need to deploy resources in multiple AWS regions.

We can write the script to create resources in different AWS accounts (here different account profiles should set).

```
provider "aws" {  
    region = "ap-south-2"  
    profile = "default"  
}  
  
provider "aws" {  
    alias = "prod"  
    region = "ap-south-1"  
    profile = "production"  
}  
  
resource "aws_s3_bucket" "dev_bucket" {  
    provider = aws  
    bucket = "ccit-dev-bucket"  
}  
  
resource "aws_s3_bucket" "prod_bucket" {  
    provider = aws.prod  
    bucket = "ccit-prod-bucket"  
}
```



Terraform

2. You need to use multiple AWS accounts.

We can write the script to create resources in different AWS accounts (here different account profiles should set).

```
provider "aws" {  
    region = "ap-south-2"  
}  
  
provider "aws" {  
    alias  = "mumbai"  
    region = "ap-south-1"  
}  
  
resource "aws_s3_bucket" "ccithyd" {  
    provider = aws  
    bucket   = "ccithydbucket123"  
}  
  
resource "aws_s3_bucket" "ccitmumbai" {  
    provider = aws.mumbai  
    bucket   = "ccitmumaibucket123"  
}
```

AWS Configure - Add default profiles
vim ~/.aws/credentials
Add different profile



Terraform

3. You want different credentials for different services.

We can create resources in different clouds (AWS, Azure, GCP etc..)

```
# AWS Provider
provider "aws" {
  region = "us-east-1"
}

# Azure Provider
provider "azurerm" {
  features {}
}

# AWS S3 Bucket
resource "aws_s3_bucket" "my_bucket" {
  bucket = "my-terraform-s3-bucket"
  acl    = "private"

  tags = {
    Name      = "MyS3Bucket"
    Environment = "Dev"
  }
}
```

```
# Azure Resource Group
resource "azurerm_resource_group" "rg" {
  name      = "MyResourceGroup"
  location  = "East US"
}

# Azure Storage Account
resource "azurerm_storage_account" "storage" {
  name           = "mystorageaccount123"
  resource_group_name = azurerm_resource_group.rg.name
  location       = azurerm_resource_group.rg.location
  account_tier   = "Standard"
  account_replication_type = "LRS"
}

# Azure Blob Storage Container
resource "azurerm_storage_container" "container" {
  name           = "mycontainer"
  storage_account_name = azurerm_storage_account.storage.name
  container_access_type = "private"
}
```



Terraform workspaces



Terraform Workspaces



- Terraform workspaces allow you to maintain multiple environments (e.g., development, testing, production)
- using the same Terraform configuration but with different states.
- Each workspace has its own state file (`terraform.tfstate.d` folder)
- resources created/managed in one workspace do not affect other workspaces.
- the default workspace in Terraform is default



Terraform Workspaces

- `terraform workspace list` : to list the workspaces.
- `terraform workspace new <<work space name>>` : to create workspace.
- `terraform workspace show` : to show current workspace.
- `terraform workspace select <<work space name>>` : to switch to <<work space name>> workspace.
- `terraform workspace delete <<work space name>>` : to delete <<work space name>> workspace.

**Dev infra
WS**

**Test infra
WS**

**Prod infra
WS**

Note

- We need to empty the workspace before delete.
- We can't delete current workspace, we can switch and delete.
- We can't delete default workspace.



Terraform

Parallelism



Terraform

Parallelism

- Terraform executes resource creation, updates, and deletions in parallel whenever possible to speed up deployment.
- However, it also respects dependencies between resources, ensuring that dependent resources are processed in the correct order.
- it will execute all the resources at a time. by default parallelism limit is 10.
- `terraform apply -auto-approve -parallelism=1`

NOTE: IT WILL APPLICABLE FOR BOTH APPLY & DESTROY.



Terraform

Dependency



Terraform

Dependency

The `depends_on` argument in Terraform is used to explicitly define dependencies between resources. This ensures that one resource is created, updated, or destroyed only after another resource has been properly handled.

- We use “`depends_on`” keyword to implement explicit dependency.

Why Use `depends_on`?

Terraform automatically determines dependencies between resources based on references.

However, in some cases, dependencies are not directly referenced, and Terraform may try to create resources in parallel.

- ✓ A resource does not have a direct reference but must still wait for another resource.
- ✓ There are implicit dependencies, like provisioning **network resources before instances**.
- ✓ A resource needs to be updated **only after** another has changed.



Terraform

Example Use Cases

1 Creating an Instance Only After a Security Group is Ready

```
resource "aws_security_group" "sg" {  
    name      = "example-sg"  
    description = "Allow web traffic"  
}  
  
resource "aws_instance" "web" {  
    ami          = "ami-09e23b3de35f110f6"  
    instance_type = "t3.micro"  
  
    depends_on = [aws_security_group.sg]  
}
```

2 Ensuring a Load Balancer is Created Before Auto Scaling

```
resource "aws_lb" "my_lb" {  
    name          = "my-load-balancer"  
    internal      = false  
    load_balancer_type = "application"  
}  
  
resource "aws_autoscaling_group" "asg" {  
    desired_capacity = 2  
    min_size        = 2  
    max_size        = 5  
    depends_on      = [aws_lb.my_lb]  
}
```

- ◆ Here, Terraform does **not** see a direct reference to `aws_security_group.sg` in `aws_instance.web`, so `depends_on` ensures Terraform creates the security group **before** launching the instance.

- ◆ Even though the ASG doesn't directly reference the load balancer, it must be created first.



⚠ When Not to Use depends_on

- ✖ If Terraform already understands the dependency from resource references, you don't need depends_on.
- ✖ Using depends_on incorrectly can make Terraform slower by forcing unnecessary sequential execution
- ✖ If Terraform already understands the dependency from resource references, you don't need depends_on.

Why Use depends_on?

- ◆ depends_on forces Terraform to respect a specific dependency
- ◆ Use it only when Terraform does not automatically detect dependencies.
- ◆ It's useful for networking, load balancers, and security groups.
- ◆ Avoid unnecessary use to keep Terraform plans efficient.

Terraform

- State file management (Securing, locking)
- Terraform lifecycle (Create before destroy, Prevent destroy, Ignore changes)
- Terraform Modules





State file management (Securing, locking)



Simple Terraform structure

```
2. root@ip-172-31-1-147:~/CCIT
provider "aws" {
  region = "ap-south-2"
}

# Create IAM user
resource "aws_iam_user" "ccit_user" {
  name = "CCITDev1"
}

# Create S3 bucket with a dynamic name
resource "aws_s3_bucket" "ccit_bucket" {
  bucket = "ccits3jan"
}

# VPC
resource "aws_vpc" "CCITVPC" {
  cidr_block          = "10.0.0.0/22"
  enable_dns_support = true
  enable_dns_hostnames = true
  tags = { Name = "CCITVPC" }
}
```

CCIT Folder





State file management (Securing, locking)

`terraform.tfstate` file



- Terraform Stores the infrastructure information on state file. it will automatically refresh when we run plan, apply & destroy.
- In Terraform, a backend is a configuration that determines how and where Terraform stores its state file & how it manages operations like apply, plan, and destroy.
- By default, Terraform uses a local backend. It stores the state file as `terraform.tfstate` in the local folder in JSON format.
- If we delete any resource, it stores information in `terraform.tfstate.backup`.



Terraform State File Locking.



Terraform State File Locking.

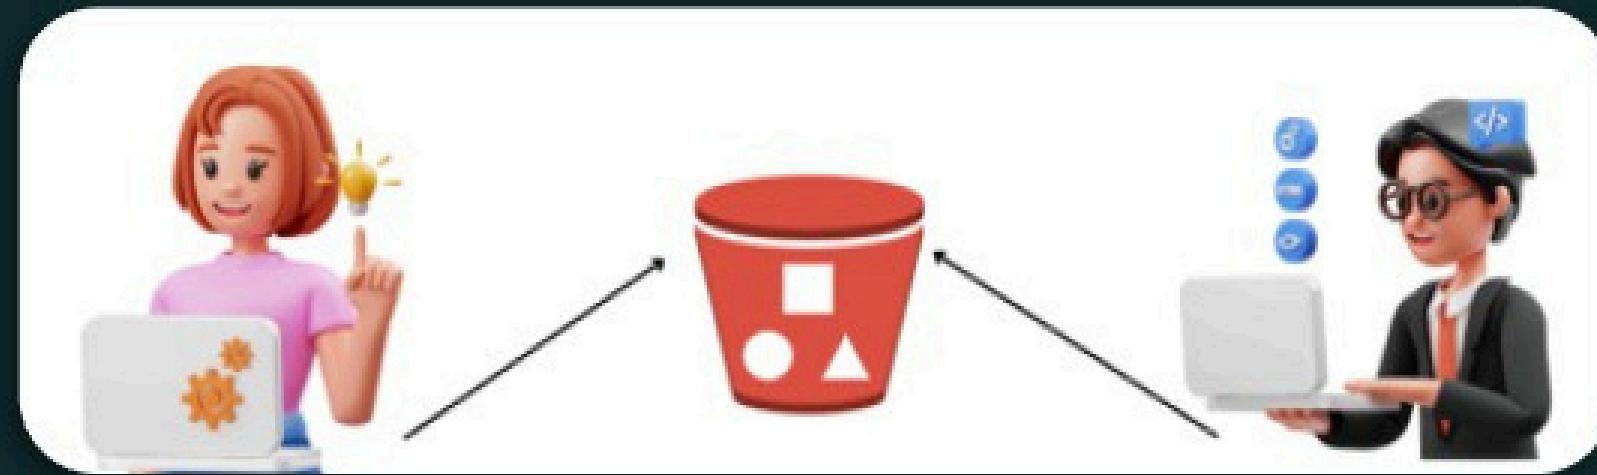
`terraform.tfstate` file



- In Real time once we complete our work, we need to lock state file.
- It ensures that only one operation can be executed at a time.
- Once you lock state file you can't modify the infrastructure anymore.
- When two people working on state file at a time it will be locked automatically.
- Unfortunately, if two people run “apply” at same time, it can lead to unpredictable results, such as creating duplicate resources or destroying the wrong infrastructure.
- Not all Terraform backends support locking.



Why locking happened?



- when 2 developers work on the same project with same state file then the locking will be happened.
- if state file is locked only first operation execute and second operation waits
- to remove state lock use: `terraform force-unlock <LOCK-ID>`
- after adding DynamoDB run: `terraform init -reconfigure`



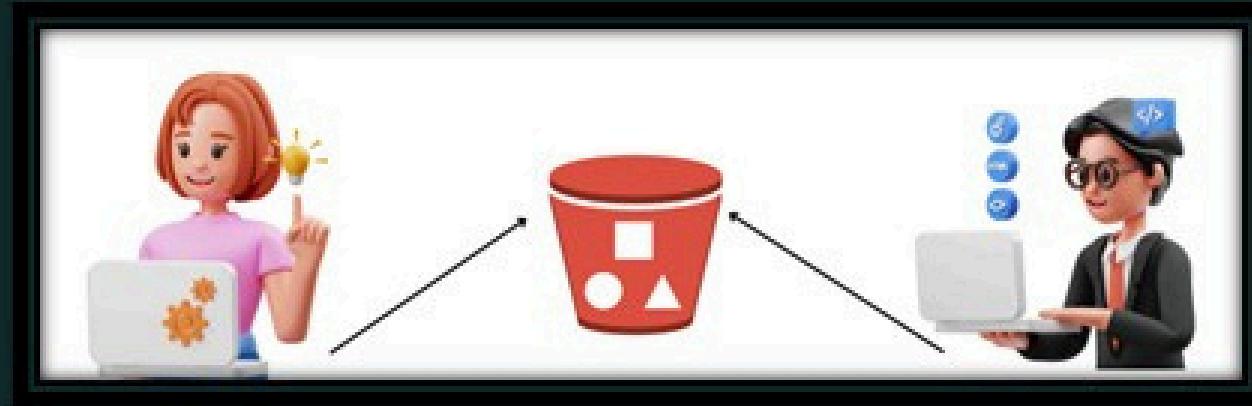
Backend	Supports Locking	Notes
Amazon S3 + DynamoDB	Yes	Requires DynamoDB table for locking.
Google Cloud Storage	Yes	Native support for state locking.
Azure Blob Storage	Yes	Native support for state locking.
HashiCorp Consul	Yes	Uses Consul's locking mechanism.
Terraform Cloud	Yes	Automatic locking with Terraform Cloud.
Alibaba Cloud OSS	Yes	Supports remote state locking.
Etcd	Yes	Distributed locking via etcd.
Local Backend	No	No support for locking (single-user only).
Git (as backend)	No	Not recommended, no locking support.



S3 backend setup for state file



S3 backend setup for state file



- Terraform used local backend to manage state file.
- But in that local backend only one person can be able to access it.
- in Real time it's often necessary to have a centralized, consistent, and secure storage mechanism for the state file
- Amazon S3 is a popular choice for this, and when combined with DynamoDB for locking, it ensures safe, consistent operations.

Advantages

- Global Access
- Team Collaboration
- Secure and Scalable
- Centralized State Management
- State File Versioning
- Disaster Recovery



S3 and DynamoDB configuration.

```
terraform {
  backend "s3" {
    bucket = "rahamterrabucketforstatefile"
    key    = "prod/terraform.tfstate"
    region = "us-east-1"
    dynamodb_table = "mytableforlock"
  }
}
```

- Add that block to existing code and run `terraform init -upgrade`
- DynamoDB -- > create table -- > Partition key: LockID -- > create
- now after apply state file will go to s3 bucket.
- dev-1 type destroy and dev-2 type apply now state file locked.
- you can check lock-id in new items of table.
- once destroy done for dev-1 state file will be unlocked and dev-2 can work.



Backup file location/ status

Backend Type	Where is tfstate stored?	How is backup handled?
Local Backend	Local directory	terraform.tfstate.backup
S3 Backend	AWS S3	Use S3 Versioning
Terraform Cloud	Terraform Cloud	Terraform manages backups



Simple Terraform structure

```
2. root@ip-172-31-1-147:~/CCIT
provider "aws" {
  region = "ap-south-2"
}

# Create IAM user
resource "aws_iam_user" "ccit_user" {
  name = "CCITDev1"
}

# Create S3 bucket with a dynamic name
resource "aws_s3_bucket" "ccit_bucket" {
  bucket = "ccits3jan"
}

# VPC
resource "aws_vpc" "CCITVPC" {
  cidr_block          = "10.0.0.0/22"
  enable_dns_support = true
  enable_dns_hostnames = true
  tags = { Name = "CCITVPC" }
}
```

CCIT Folder





Terraform

Terraform Modules



Simple Terraform structure

```
provider "aws" {
  region = "ap-south-2"
}

# Create IAM user
resource "aws_iam_user" "ccit_user" {
  name = "CCITDev1"
}

# Create S3 bucket with a dynamic name
resource "aws_s3_bucket" "ccit_bucket" {
  bucket = "ccits3jan"
}

# VPC
resource "aws_vpc" "CCITVPC" {
  cidr_block          = "10.0.0.0/22"
  enable_dns_support = true
  enable_dns_hostnames = true
  tags = { Name = "CCITVPC" }
}
```

```
provider "aws" {
  region = "ap-south-2"
}

# Create IAM user
module "iam_user" {
  source    = "./modules/iam"
}

# Create S3 bucket
module "s3_bucket" {
  source    = "./modules/s3"
}

# Create VPC
module "vpc" {
  source    = "./modules/vpc"
}

module "vpc" {
  source  = "terraform-aws-modules/vpc/aws"
  version = "5.19.0"
}
```

```
resource "aws_iam_user" "user" {
  name = "CCITDev1"
}
```

```
resource "aws_s3_bucket" "bucket" {
  bucket = "ccits3jan"
}
```

```
resource "aws_vpc" "this" {
  cidr_block          = "10.0.0.0/22"
  enable_dns_support = true
  enable_dns_hostnames = true
  tags = {
    Name = "CCITVPC"
  }
}
```

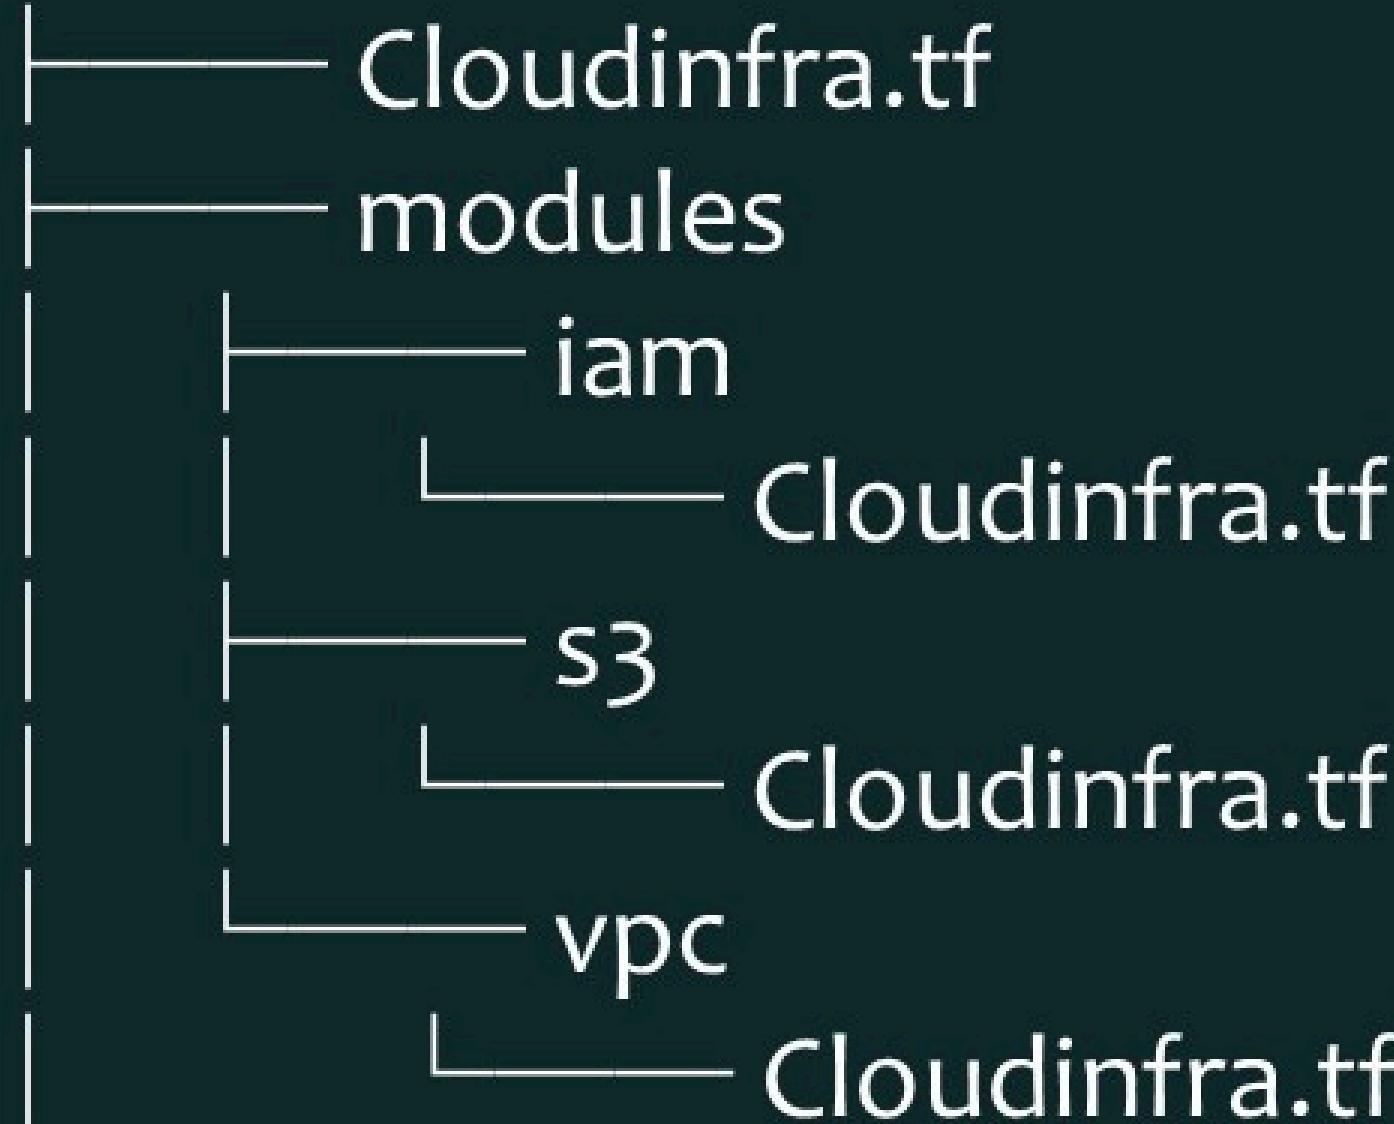


Modules in Terraform

- It divides the code into folder structure.
- Modules are group of multiple resources that are used together.
- This makes your code easier to read and reusable across your organization.
- we can publish modules for others to use.
- each module will be having separate plugin
- modules plugins will be stored on .terraform/modules/



Modules in Terraform



```
provider "aws" {
  region = "ap-south-2"
}

# Create IAM user
module "iam_user" {
  source    = "./modules/iam"
}

# Create S3 bucket
module "s3_bucket" {
  source    = "./modules/s3"
}

# Create VPC
module "vpc" {
  source    = "./modules/vpc"
}

module "vpc" {
  source  = "terraform-aws-modules/vpc/aws"
  version = "5.19.0"
}
```

```
resource "aws_iam_user" "user" {
  name = "CCITDev1"
}
```

```
resource "aws_s3_bucket" "bucket" {
  bucket = "ccits3jan"
}
```

```
resource "aws_vpc" "this" {
  cidr_block          = "10.0.0.0/22"
  enable_dns_support = true
  enable_dns_hostnames = true

  tags = {
    Name = "CCITVPC"
  }
}
```



Types Modules

```
provider "aws" {
  region = "ap-south-2"
}

# Create IAM user
module "iam_user" {
  source  = "./modules/iam"
}

# Create S3 bucket
module "s3_bucket" {
  source    = "./modules/s3"
}

# Create VPC
module "vpc" {
  source  = "./modules/vpc"
}

module "vpc" {
  source  = "terraform-aws-modules/vpc/aws"
  version = "5.19.0"
}
```

```
resource "aws_iam_user" "user" {
  name = "CCITDev1"
}
```

```
resource "aws_s3_bucket" "bucket" {
  bucket = "ccits3jan"
}
```

```
resource "aws_vpc" "this" {
  cidr_block          = "10.0.0.0/22"
  enable_dns_support = true
  enable_dns_hostnames = true

  tags = {
    Name = "CCITVPC"
  }
}
```

Root Module: This is the main directory where Terraform commands are run. All Terraform configurations belong to the root module.

Child Modules: These modules are called by other modules.



Module Sources

- Modules can be sourced from a number of different locations, including both local and remote sources.
- we can get modules form Terraform Registry, HTTP URLs, S3 buckets & Local.

```
provider "aws" {
region = "us-east-1"
}

module "vpc" {
  source  = "terraform-aws-modules/vpc/aws"
  version = "5.13.0"
}
```



Terraform

Terraform life cycle



Terraform life cycle

Terraform Lifecycle: Managing Resource Behavior 🚀

Terraform provides a **lifecycle** block to control how resources are created, updated, and deleted. It helps prevent accidental destruction, force resource recreation, and ignore specific changes.

Meta-Argument	Description
<code>create_before_destroy</code>	Ensures a new resource is created before the old one is deleted.
<code>prevent_destroy</code>	Prevents Terraform from accidentally deleting a resource.
<code>ignore_changes</code>	Ignores specific attribute changes to avoid unnecessary updates.



Terraform life cycle

`create_before_destroy`

- ✓ Ensures a new resource is created before deleting the old one.
 - ◆ Use case: When replacing a resource without downtime.

```
resource "aws_instance" "example"
{
    ami = "ami-12345678"
    instance_type = "t2.micro"

    lifecycle
    {
        create_before_destroy = true
    }
}
```

Behavior:

1. Terraform creates the new instance.
2. Once created, Terraform destroys the old one.
3. Prevents downtime since there is no gap between deletion and recreation.



Terraform life cycle

prevent_destroy

- ✓ Prevents accidental deletion of critical resources.
- ◆ **Use case:** Protecting **databases, S3 buckets, or production VMs**.

```
resource "aws_instance" "example"
{
    ami = "ami-12345678"
    instance_type = "t2.micro"

    lifecycle
    {
        prevent_destroy = true
    }
}
```

Behavior:

1. Terraform doesn't allow to destroy the instance



Terraform life cycle

ignore_changes

- ✓ Ignores changes to specific attributes.

- ◆ **Use case:** When an attribute changes outside Terraform (e.g., manual AWS console updates).

```
resource "aws_instance" "example"
{
  ami = "ami-12345678"
  instance_type = "t2.micro"

  lifecycle
  {
    ignore_changes = [tags]
  }
  tags = {
    Name = "Example-EC2"
  }
}
```

Behavior:

1. Terraform refuses the changes mentioned.

Terraform

- For each loop in Terraform
- Terraform locals
- Dynamic block
- Built-In functions
- Blocks, args & expressions
- Advantages & disadvantages





Terraform

For each loop in Terraform



For each loop in Terraform

- The "for_each" in Terraform is used to loop over a map or a set to create multiple resources dynamically.
- For_each syntax:

```
resource "resource_type" "name" {  
    for_each = <map or set>  
  
    attribute = each.value # Access the value  
    another_attribute = each.key # Access the key (only for maps)  
}
```



Terraform

List in Terraform



Terraform

List in Terraform

- A list in Terraform is an ordered collection of elements where each item has a specific position (index). Duplicates are allowed.
- Syntax

```
variable "my_list" {  
    type    = list(string)  
    default = ["bucket-1", "bucket-2", "bucket-3"]  
}
```

Set in Terraform

- A set in Terraform is an unordered collection of unique elements. Duplicates are not allowed.
- Syntax

```
variable "my_set" {  
    type    = set(string)  
    default = ["bucket-1", "bucket-2", "bucket-3"]  
}
```



Terraform

Difference between List and Set



Terraform

Key Differences Between Lists and Sets

Feature	List ✓	Set ✓
Maintains order	✓ Yes	✗ No
Allows duplicates	✓ Yes	✗ No
Uses index (e.g., var.my_list[0])	✓ Yes	✗ No
Used with count	✓ Yes	✗ No
Used with for_each	✗ No	✓ Yes



Terraform

Map



Terraform

Map

Used to pass both key and value to variables.

KEY-VALUE can also called as object or Dictionary.

```
variable "variable_name" {  
    type  = map(string)  
    default = {  
        key1 = "value1"  
        key2 = "value2"  
    }  
}
```



Terraform

Terraform locals



Terraform

Terraform locals:

In Terraform, locals are used to define local values within a module. These local values help simplify configurations by assigning names to expressions, making them easier to reference throughout the configuration.

Why Use locals?

- 1. Improve Readability** – Instead of repeating complex expressions, you can define them once and use them multiple times.
- 2. Avoid Repetition** – Helps in maintaining the DRY (Don't Repeat Yourself) principle.
- 3. Enhance Maintainability** – If a value needs to be updated, you only change it in one place.

```
# Define Local Variables
locals {
  bucket_name = "ccitmarch25bucket"
  instance_type = "t3.micro"
  ami_id      = "ami-0150f9c33a6958e02"
}
```

```
# Create S3 Bucket
resource "aws_s3_bucket" "my_bucket" {
  bucket = local.bucket_name
  acl    = "private"

  tags = {
    Name = "MyS3Bucket"
  }
}
```

```
# Launch EC2 Instance
resource "aws_instance" "my_ec2" {
  ami           = local.ami_id
  instance_type = local.instance_type

  tags = {
    Name = "MyEC2Instance"
  }
}
```



Terraform

Terraform locals:

Feature	variable (input variable)	local (local value)
Purpose	Accept external input (from user, CLI, tfvars)	Define internal reusable values
Can be overridden?	<input checked="" type="checkbox"/> Yes — via CLI, tfvars, etc.	<input type="checkbox"/> No — fixed at definition time
Use case	Configure environment, region, name, etc.	Simplify expressions, avoid duplication
Where used	Module input	Module internal only
Change from outside?	<input checked="" type="checkbox"/> Yes	<input type="checkbox"/> No

Summary:

Use variable when:	Use local when:
You expect input from outside	You want to clean up internal logic
Config needs to be flexible	Computed values, reused expressions
Used across modules	Used only inside one module/file



Terraform

Dynamic block



Dynamic block

Dynamic blocks in Terraform are a powerful feature that allow for the dynamic generation of nested blocks within resources, data sources, providers, or provisioners. Here's a concise overview:

- 1. Purpose:** Dynamic blocks enable the creation of multiple nested configurations based on variable input, reducing code repetition and adhering to the DRY (Don't Repeat Yourself) principle.
- 2. Structure:** A dynamic block consists of a `for_each` argument to iterate over a collection and a content block that defines the structure to be dynamically generated.
- 3. Common Use Cases:** Dynamic blocks are particularly useful for resources that require multiple nested blocks, such as defining multiple ingress rules in a security group or configuring multiple subnets within a VPC.
- 4. Advantages:** They enhance code maintainability by reducing redundancy and allow for more flexible and scalable configurations.
- 5. Limitations:** Overusing dynamic blocks can make configurations harder to read and maintain. It's essential to balance their use to avoid unnecessary complexity.



Terraform

Terraform Built-in functions



Terraform Built-in functions

Terraform provides a variety of **built-in functions** to manipulate and transform values. These functions are mainly used in expressions and variables.

1 String Functions 

2 Numeric Functions 

3 Collection Functions (Lists & Maps) 

4 IP & CIDR Functions 

5 Type Conversion Functions 

6 Encoding & Hashing Functions 

7 File & Path Functions 



1 String Functions abc

Function	Description	Example
upper()	Converts a string to uppercase	upper("terraform") → "TERRAFORM"
lower()	Converts a string to lowercase	lower("TERRAFORM") → "terraform"
length()	Returns length of a string	length("Terraform") → 9
substr()	Extracts part of a string	substr("Terraform", 0, 4) → "Terr"
replace()	Replaces substrings	replace("hello world", "world", "Terraform") → "hello Terraform"
trimspace()	Removes spaces from start and end	trimspace(" Terraform ") → "Terraform"
split()	Splits a string into a list	split("", "a,b,c") → ["a", "b", "c"]
join()	Joins a list into a string	join("-", ["a", "b", "c"]) → "a-b-c"

2 Numeric Functions 1234

Function	Description	Example
abs()	Absolute value	abs(-10) → 10
ceil()	Rounds up	ceil(4.3) → 5
floor()	Rounds down	floor(4.9) → 4
max()	Returns the largest number	max(10, 20, 30) → 30
min()	Returns the smallest number	min(10, 20, 30) → 10
pow()	Exponentiation	pow(2, 3) → 8
log()	Logarithm	log(100, 10) → 2



Terraform

3 Collection Functions (Lists & Maps)

Function	Description	Example
length()	Returns number of elements	length(["a", "b", "c"]) → 3
element()	Gets an element by index	element(["a", "b", "c"], 1) → "b"
concat()	Combines lists	concat(["a"], ["b", "c"]) → ["a", "b", "c"]
contains()	Checks if a list has an item	contains(["a", "b"], "b") → true
flatten()	Flattens nested lists	flatten([["a", "b"], ["c"]]) → ["a", "b", "c"]
keys()	Returns keys of a map	keys({name="John", age=30}) → ["name", "age"]
values()	Returns values of a map	values({name="John", age=30}) → ["John", 30]
merge()	Merges multiple maps	merge({a=1}, {b=2}) → {a=1, b=2}

4 IP & CIDR Functions

Function	Description	Example
cidrsubnet()	Creates a subnet from a CIDR	cidrsubnet("192.168.1.0/24", 2, 1) → "192.168.1.64/26"
cidrhost()	Gets an IP from a CIDR	cidrhost("192.168.1.0/24", 5) → "192.168.1.5"
cidrnetmask()	Gets the netmask length	cidrnetmask("192.168.1.0/24") → "255.255.255.0"

5 Type Conversion Functions

Function	Description	Example
tostring()	Converts to string	tostring(123) → "123"
tonumber()	Converts to number	tonumber("42") → 42
tobool()	Converts to boolean	tobool("true") → true



Terraform

6 Encoding & Hashing Functions 🔑

Function	Description	Example
base64encode()	Encodes to Base64	base64encode("hello") → "aGVsbG8="
base64decode()	Decodes Base64	base64decode("aGVsbG8=") → "hello"
md5()	Generates MD5 hash	md5("hello") → "5d41402abc4b2a76b9719d911017c592"
sha256()	Generates SHA-256 hash	sha256("hello") → "2cf24dba5fb0a30e26e83b2ac5b9e29e1b161e5c1fa7425e73043362938b9824"

7 File & Path Functions 📁

Function	Description	Example
file()	Reads file content	file("config.json")
dirname()	Gets the directory name	dirname("/usr/local/bin") → "/usr/local"
basename()	Gets the filename	basename("/usr/local/bin/script.sh") → "script.sh"
path.cwd	Current working directory	"\${path.cwd}"



Terraform

Blocks, args & expressions



Blocks

In Terraform, configurations are written using **blocks**, **arguments**, and **expressions**.

1 Blocks (Structural Units)

A block is a container for configuration settings.

Each block has:

1. A block type (e.g., resource, provider, variable)
2. A label (optional, depends on the block type)
3. Arguments inside {} brackets

Block Type	Purpose	Usually Covered In
resource	Creates infrastructure	<input checked="" type="checkbox"/> Main list
provider	Configures cloud/API access	<input checked="" type="checkbox"/> Main list
module	Reuses other Terraform code	<input checked="" type="checkbox"/> Main list
variable	Accepts input values	<input checked="" type="checkbox"/> Main list
locals	Defines reusable internal values	<input checked="" type="checkbox"/> Main list
output	Displays/export values after apply	<input checked="" type="checkbox"/> Main list
data	Fetches existing cloud resources	<input checked="" type="checkbox"/> Main list
terraform	Sets backend, version, providers	<input checked="" type="checkbox"/> Setup/meta config



Blocks

In Terraform, configurations are written using **blocks**, **arguments**, and **expressions**.

1 Blocks (Structural Units)

A block is a container for configuration settings.

Each block has:

1. A block type (e.g., resource, provider, variable)
2. A label (optional, depends on the block type)
3. Arguments inside {} brackets

Provider Block

```
provider "aws" {  
  region = "ap-south-2"  
}
```

Local Block

```
locals {  
  bucket_name  = "ccitmarch25bucket"  
  instance_type = "t3.micro"  
  ami_id       = "ami-0150f9c3a6958e02"  
}
```

Resource block

```
resource "aws_s3_bucket" "my_bucket" {  
  bucket = "ccitmar25bucket"
```

Variable block

```
variable "instance_type" {  
  default = "t2.micro"  
}
```



2 Arguments (Key-Value Pairs Inside Blocks)

An argument is a key-value pair inside a block that defines setting

👉 Arguments can accept:

- String** ("t3.micro")
- Number** (10)
- Boolean** (true / false)
- List** ("dev", "test")
- Map** ({ name = "server", env = "dev" })

```
resource "aws_instance" "inst" {  
    ami          = "ami-0c55b1" # ami is an argument  
    instance_type = "t3.micro" # ins_type is an argument  
}
```



Terraform

Expressions

Expressions compute values dynamically based on variables, functions, or resource attributes.

Types of Expressions

- Variables Expressions
- Function Expressions
- Resource Attributes
- Conditional Expressions

Variable expressions

```
variable "inst_type" {
  default = "t2.micro"
}

resource "aws_instance" "inst" {
  instance_type = var.inst_type # Using a variable in an expression
}
```

Function Expressions

```
locals {
  upper_name = upper("terraform") # Converts to uppercase
}
```

Resource Attributes

```
output "public_ip" {
  value = aws_instance.example.public_ip # Fetching EC2 public IP dynamically
}
```

Conditional Expressions

```
locals {
  instance_type = var.is_production ? "t3.large" : "t2.micro"
}
```



Terraform

Advantages & disadvantages



Terraform

✓ Advantages of Terraform

1 Multi-Cloud Support

- Works with AWS, Azure, GCP, Kubernetes, etc.
- Allows managing resources across multiple clouds.

2 Declarative Language (HCL - HashiCorp Configuration Language)

- Defines what the infrastructure should look like rather than how to achieve it.
- Easier to read and manage.

3 Modular & Reusable Code

- Supports modules to create reusable components.
- Example: A single VPC module can be used across different projects.

4 Supports Infrastructure Automation

Integrates with CI/CD pipelines (Jenkins, GitHub Actions, etc.).

5 Open-Source & Extensible

- Free to use, with Terraform Cloud offering extra features.
- Can be extended with providers & plugin



✖ Disadvantages of Terraform

1 State File Management Issues

- The `.tfstate` file contains sensitive information (e.g., secrets, credentials).
- Needs remote storage (S3, Terraform Cloud) for security and collaboration.

2 Learning Curve

- HCL syntax is different from traditional scripting languages.
- Troubleshooting Terraform errors can be complex.

3 Longer Execution Time for Large Deployments

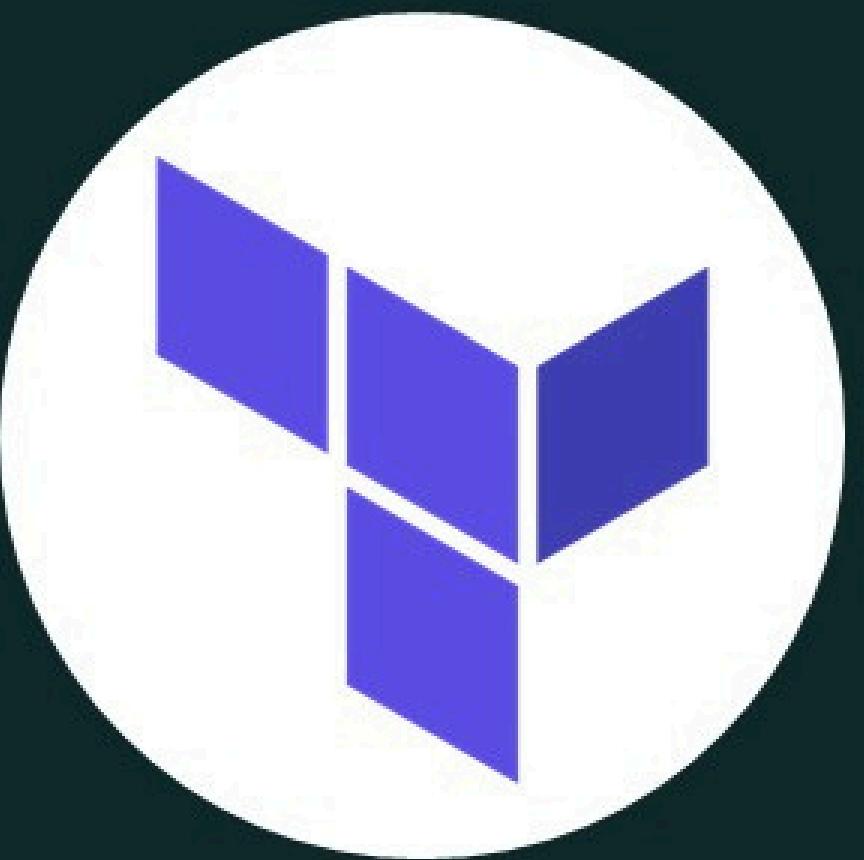
- Terraform re-evaluates dependencies before each apply, which can be slow for large infrastructures.
- Performance may be impacted when managing thousands of resources.

5 Destroy Command is Dangerous

- Running `terraform destroy` removes all infrastructure.
- Needs extra safety measures to avoid accidental deletions.

Terraform

- Terraform CLI
- Providers block
- Local version changing
- Provisioners (File, Local-exec, Remote exec)
- Variable precedence





Terraform

Terraform CLI



Terraform

🔧 Initialization and Setup

Command	Description
terraform init	Initialize Terraform in a directory
terraform version	Show Terraform version
terraform providers	List providers used in the configuration

📁 Workspaces

Command	Description
terraform workspace list	List all workspaces
terraform workspace show	Show current workspace
terraform workspace new	Create a new
terraform workspace select	Switch to a different
terraform workspace delete	Delete a workspace

⚙️ Core Workflow

Command	Description
terraform plan	Show execution plan
terraform apply	Apply changes
terraform destroy	Destroy managed infrastructure
terraform validate	Validate configuration syntax and structure
terraform fmt	Format configuration files
terraform show	Show current or planned state
terraform output	Show output variables

⌚ Resource Lifecycle

Command	Description
terraform taint	Mark a resource for recreation
terraform untaint	Remove taint from a resource
terraform import	Import existing infrastructure into Terraform
terraform refresh	Update the state file with real-world resources



Terraform

💡 Debugging and Logs

Command	Description
<code>terraform console</code>	Interactive console for Terraform expressions
<code>terraform debug</code>	Start a debug session (advanced use)
<code>TF_LOG=DEBUG</code> (env var)	Enable debug logging for troubleshooting

📦 Modules & Providers

Command	Description
<code>terraform providers</code>	Show required and installed providers
<code>terraform providers lock</code>	Generate dependency lock file
<code>terraform providers mirror</code>	Download provider plugins to a local directory
<code>terraform get</code>	Download modules (for older workflows)



Terraform

Providers block



Providers block

- By default, provider plugins in Terraform change version every few weeks.
- When we run the init command, it downloads the latest plugins always.
- Some code will not work with old plugins, so we need to update them.
- To get the latest provider plugins: <https://registry.terraform.io/browse/providers>.
- When you add a new provider, “terraform init” is a must.
- `terraform providers`: Lists the providers required to run the code.
- To create infrastructure on any cloud using Terraform, all we need is a provider.

Types:

- OFFICIAL: Managed by Terraform.
- PARTNER: Managed by a third-party company.
- COMMUNITY: Managed by individuals.



Terraform

Local version changing



Terraform

Local version changing

- When you add/ update the provider, terraform init (-upgrade) is a must.
- `terraform providers` command gives you the required providers information.
- `cat .terraform.lock.hcl` : Here you can get the initiated providers information.

```
terraform {  
  required_providers {  
    aws = {  
      source = "hashicorp/aws"  
      version = "5.91.0"  
    }  
  }  
}
```

```
terraform {  
  required_providers {  
    aws = {  
      source = "hashicorp/aws"  
      version = " $\leq$ 5.91.0"  
    }  
  }  
}
```

```
terraform {  
  required_providers {  
    aws = {  
      source = "hashicorp/aws"  
      version = " $\geq$ 5.91.0"  
    }  
  }  
}
```



Update the providers

- When you add/ update the provider, `terraform init (-upgrade)` is a must.
- `terraform providers` command gives you the required providers information.
- `cat .terraform.lock.hcl` : Here you can get the initiated providers information.

Operator	Description
<code>= 1.8.0</code>	Only version 1.8.0 is allowed
<code>!= 1.8.0</code>	Any version except 1.8.0 is allowed
<code>> 1.8.0</code>	Any version greater than 1.8.0
<code>< 1.8.0</code>	Any version less than 1.8.0
<code>>= 1.8.0</code>	1.8.0 or newer
<code><= 1.8.0</code>	1.8.0 or older
<code>~> 1.8.0</code>	1.8.x versions only (e.g., 1.8.1, 1.8.5, but not 1.9.0)
<code>~> 1.8</code>	1.x versions only (e.g., 1.8.1, 1.9.0, but not 2.0.0)

```
terraform {
  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = ">=5.67.0"
    }

    azurerm = {
      source  = "hashicorp/azurerm"
      version = ">=4.24.0"
    }

    google = {
      source  = "hashicorp/google"
      version = ">=6.26.0"
    }
  }
}
```



Terraform

Terraform block and Provider Version Constraints

You can define version constraints for providers using the `required_providers` block.

It usually includes details such as required providers, backend configuration, and version constraints.

- `terraform -v`
- `terraform -version`
- `terraform --version`

```
hcl
terraform {
  required_version = ">= 1.8.0"

  required_providers {
    aws = {
      source  = "hashicorp/aws"
      version = "~> 5.91.0" # Allows 5.91.x versions but not 5.92.0
    }
  }
}
```

Setting	Description
<code>required_version = ">= 1.8.0"</code>	Ensures Terraform is version 1.8.0 or newer
<code>required_providers</code>	Specifies required providers (e.g., AWS)
<code>source = "hashicorp/aws"</code>	Fetches the AWS provider from HashiCorp's registry
<code>version = "~> 5.91.0"</code>	Allows 5.91.x versions but prevents upgrading to 5.92.0+



Provisioners (File, Local-exec, Remote exec)



🔧 What Are Provisioners?

Provisioners are used to execute scripts or commands on a local or remote machine after the resource is created.

They're often used to:

- Copy files
- Run installation commands
- Trigger configuration scripts

⚠️ **Note:** Provisioners are not the best practice for production — try to use cloud-init or user_data instead. But they're great for learning and small tasks.



Terraform

1 File Provisioner

- Transfers files from your local machine to the remote instance.
- Works only for resources with SSH access (like EC2).

 ccit.txt must exist in the same directory.

Example: Upload a Script to an EC2 Instance

```
provider "aws" {}

resource "aws_instance" "example" {
    ami          = "ami-0af13678ce46c8851"
    instance_type = "t2.micro"
    key_name      = "my-key"

    provisioner "file" {
        source      = "hello.txt"
        destination = "/home/ec2-user/ccit.txt"
    }

    connection {
        type      = "ssh"
        user      = "ec2-user"
        private_key = file("~/ssh/my-key.pem")
        host      = self.public_ip
    }
}
```



2 Local-Exec Provisioner

- Runs a command on your local machine (where Terraform is executed).
- Useful for tasks like sending notifications or running local scripts.

Example: Run a Local Script After resource Creation

```
provider "aws" {  
    region = "ap-south-2"  
}  
resource "aws_s3_bucket" "bucket" {  
    bucket = "ccits3bucketmarch"  
}  
resource "null_resource" "upload_file" {  
    depends_on = [aws_s3_bucket.bucket]  
    provisioner "local-exec" {  
        command = "aws s3 cp myfile.txt s3://ccits3bucketmarch/"  
    }  
}
```

3 Remote-Exec Provisioner

- Runs **commands on the remote machine** after SSH access.
- Used for configuration tasks like installing packages.

Example: Install Apache on EC2

```
resource "aws_instance" "example" {  
    ami          = "ami-0c55b159cbfafe1f0"  
    instance_type = "t2.micro"  
  
    connection {  
        type      = "ssh"  
        user      = "ec2-user"  
        private_key = file("~/ssh/id_rsa")  
        host      = self.public_ip  
    }  
  
    provisioner "remote-exec" {  
        inline = [  
            "sudo yum update -y",  
            "sudo yum install -y httpd",  
            "sudo systemctl start httpd"  
        ]  
    }  
}
```



Terraform

Variable precedence



Terraform

VARIABLE PRECEDENCE:

- Defines the Level Priority for Variables in Terraform
- Terraform will pick the variables in the following order:
 - CLI
 - `terraform.tfvars`
 - Environment Variable
 - `variable.tf`

Environment Variable Example:

- `export TF_VAR_instance_type="t2.medium"`
- `$env:TF_VAR_instance_type = "t2.medium"`

Terraform

- HCP Cloud introduction and features
- GitHub Communication
- Build infra with HCP





Terraform

HCP Cloud introduction and features



HCP Cloud - Introduction

- HCP means HashiCorp Cloud Platform
- It is a managed platform to automate cloud infrastructure.
- It provides privacy, security, and isolation.
- It supports multiple providers like AWS, Azure, and Google Cloud.
- It offers a suite of open-source tools for managing infrastructure, including Terraform, Vault, Consul, and Nomad.
- We can use different code repositories for a project.
- We can use variable sets to apply the same variables to all the workspaces.



Terraform

Main products

- Terraform: to create infrastructure.
- Vault: to manage the secrets and protecting sensitive data.
- Nomad: to schedule, deploy and manage applications.
- Consul: to secure service-to-service communication and networking.
- Packer: to create images for launching servers.

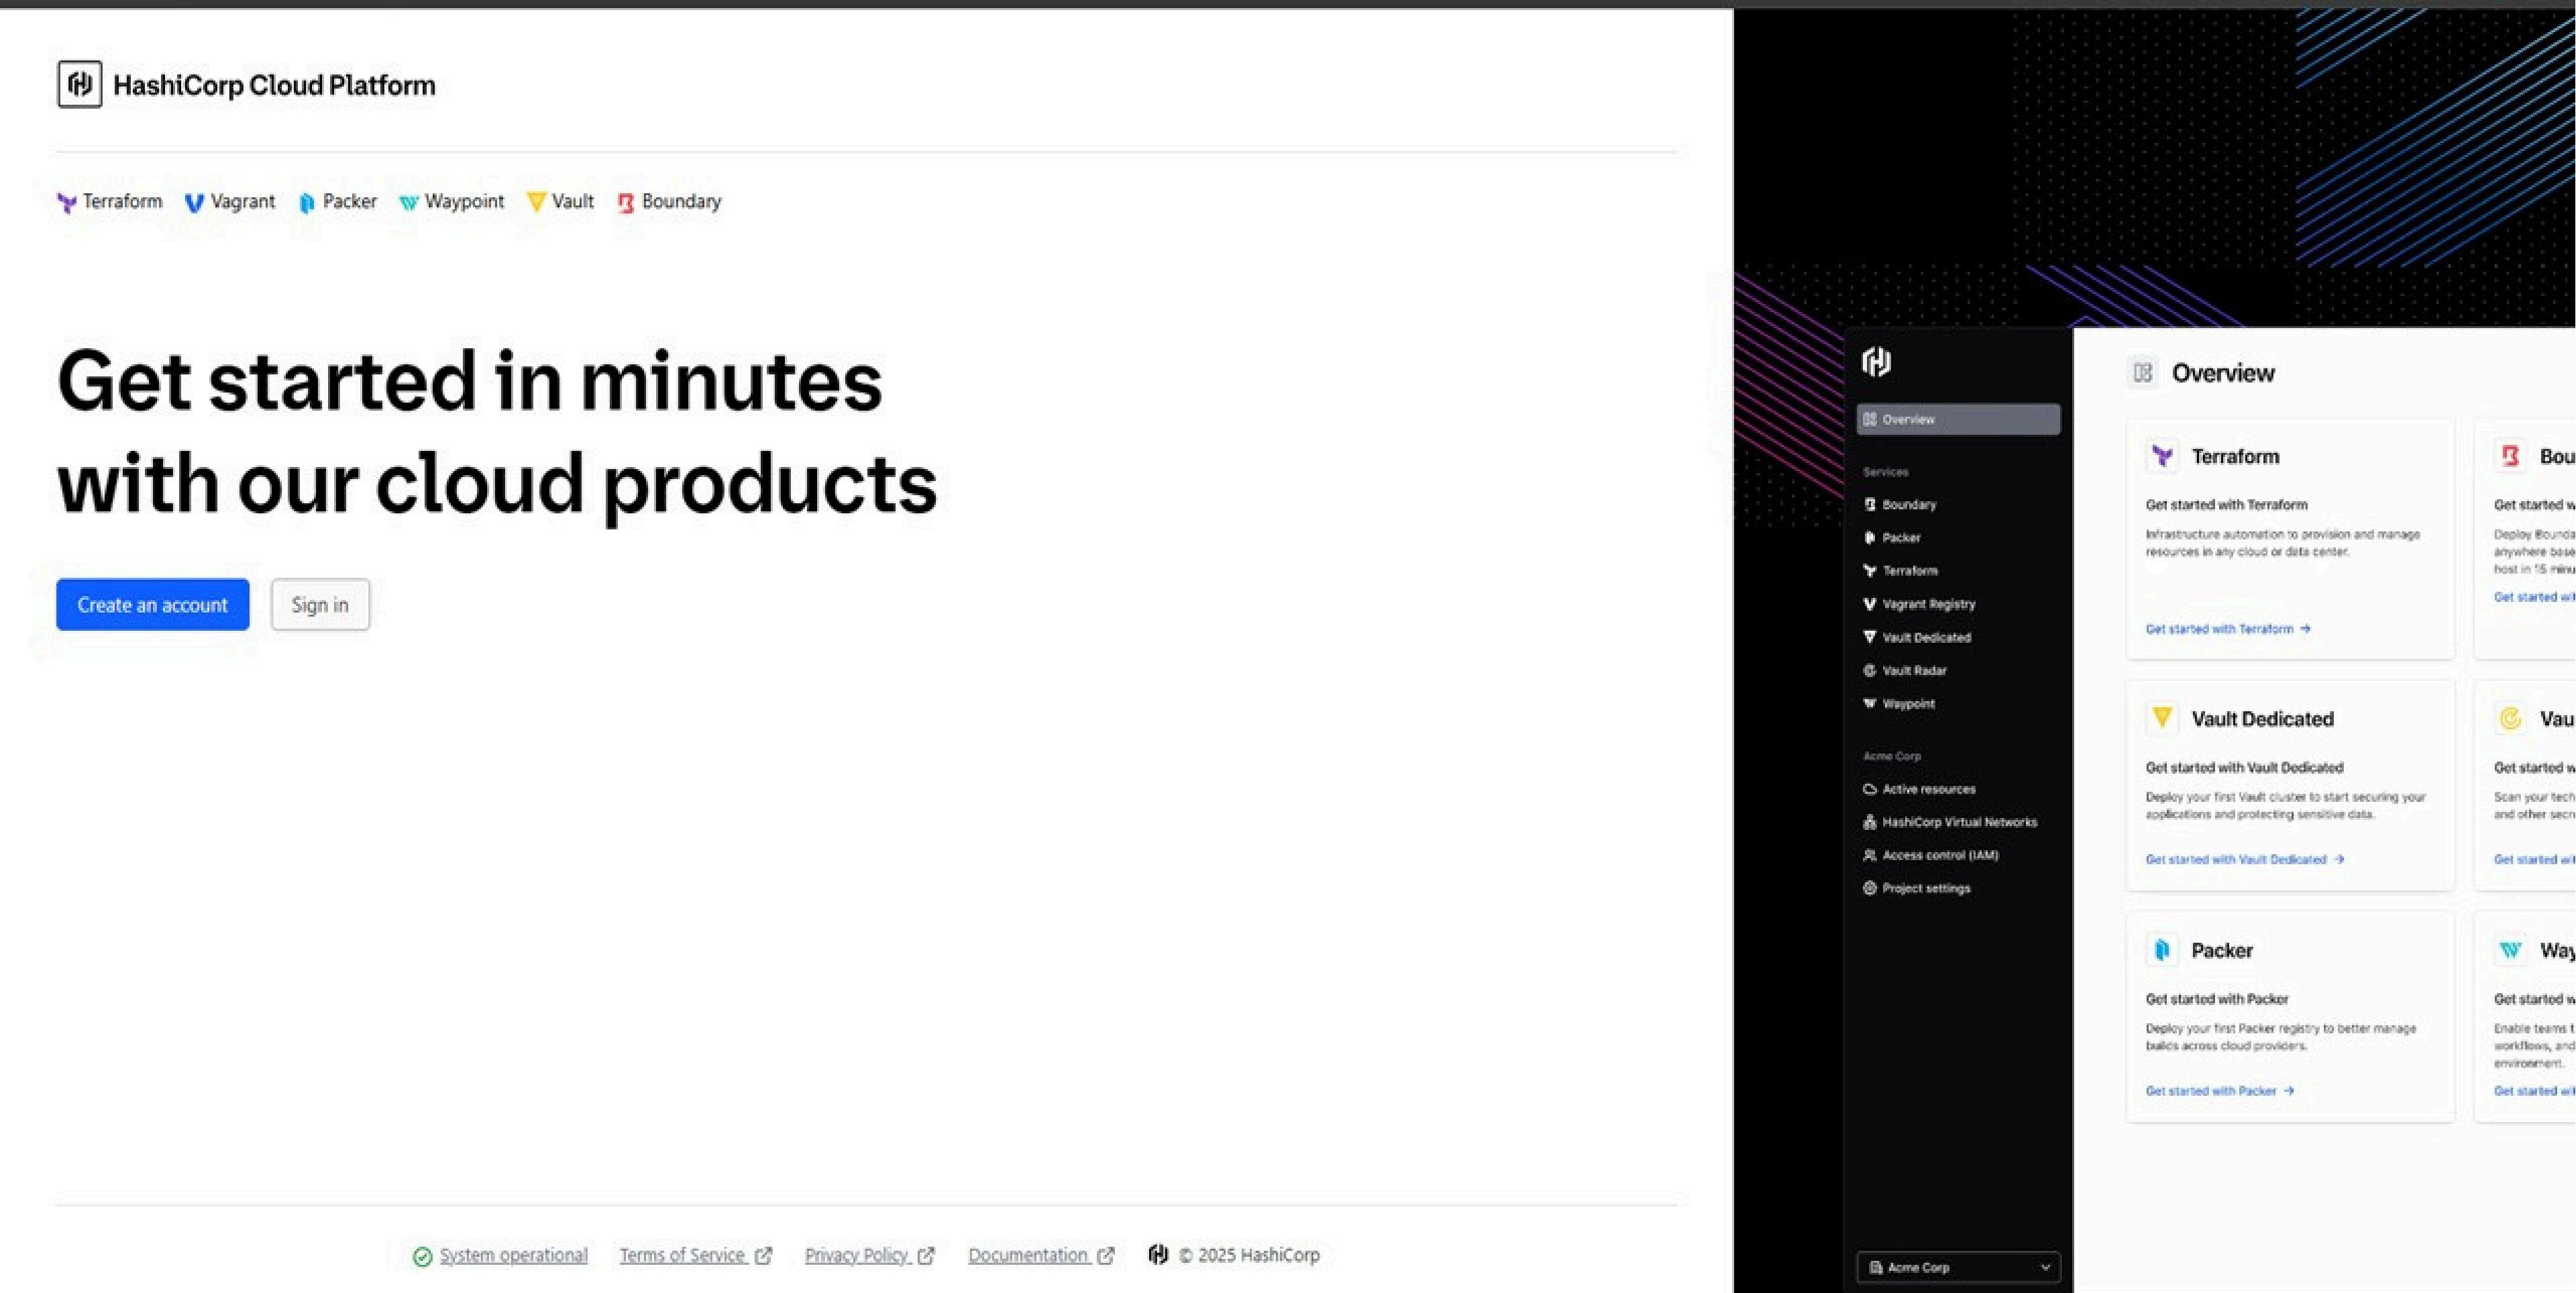
Account Creation

- Go to Google & type: HCP Cloud account Sign-In/ Sign-Up
- Email & Password
- Verify the email → Create ORG
- Click on Terraform.
- Continue with HCP account
- Create A GitHub account.
- create repo → name → add new file → write terraform code → commit



Terraform

HCP Cloud signup





Create HashiCorp Account

Enter your email or GitHub to continue

Email address* —

By continuing, you acknowledge HashiCorp's [Terms of use](#) and the [Privacy policy](#).

[Continue](#)

Already have an account? [Log in](#)

OR

 [Continue with GitHub](#)

WhatsApp

Inbox (14) - cloudtechintelugu@gmail.com

https://mail.google.com/mail/u/1/#inbox

Gmail

Compose

Inbox 14

Starred

Snoozed

Sent

Drafts

More

Labels +

Get additional protection against phishing

Turn on Enhanced Safe Browsing to get additional protection against dangerous emails

Continue No thanks

Primary Promotions 3 new Zoom — Expand your reach Social

Get started with Gmail

Customize your inbox Change profile image Import contacts and mail Get Gmail for mobile

HashiCorp Cloud Pla. Verify your email - Please verify your email address Press the button below to verify cloudtechintelugu@gmail.com for use with your HashiCorp account. Verify Or verify using this link... 3:45 PM

HashiCorp Cloud Pla. Invitation to enroll in Multifactor Authentication - Multifactor Authentication is now available MFA is an optional feature that will increase the security of your account. Enroll in Multifact... 3:45 PM

Amazon Web Services Amazon Web Services GST Invoice Available - Greetings from Amazon Web Services, We are writing to notify you that your AWS Account has GST Invoice(s) available. They can be down... Aug 2

GitHub [GitHub] A fine-grained personal access token has been added to your account - @cloudtechintelugu, a personal access token was created on your account. Hey cloudtechintelugu! ... Jul 31

GitHub [GitHub] A first-party GitHub OAuth application has been added to your account - Hey cloudtechintelugu! A first-party GitHub OAuth application (Git Credential Manager) with gist, r... Jul 31

GitHub [GitHub] A fine-grained personal access token has been added to your account - @cloudtechintelugu, a personal access token was created on your account. Hey cloudtechintelugu! ... Jul 26

GitHub [GitHub] Your GitHub launch code - Here's your GitHub launch code! Continue signing up for GitHub by entering the code below: 12527076 Open GitHub Not able to enter the code? Paste the foll... Jul 26

Zoom New Invoice Notification - Thanks for your business! Please see the attached copy of your invoice. Please refer to the payment instructions on the attached invoice. Zoom Logo Header In... Jul 24

8a129b6a9836d...

Zoom Welcome to Zoom! - Jul 24

INV314735435...

Zoom Communications. Payment Processing - Your payment is processing. Zoom Logo Header Payment processing Dear Cloud Computing In Telugu, Your recent payment of ₹1623.68 Indian Rupee for Account N... Jul 24

Document processed - Your document has been successfully processed and saved to your account. Zoom I am Wonder Document processed Dear Cloud Computing In Telugu, Thank you for Jul 24

Upgrade →

WhatsApp Verify your email - cloudtechintellugu@gmail.com +

https://mail.google.com/mail/u/1/#inbox/PMfcgzQbgcSmdqvKhtDwvXWJMrptTh

Gmail Search mail

Compose

Inbox 13 Starred Snoozed Sent Drafts More Labels +

Verify your email Inbox

HashiCorp Cloud Platform <no-reply@activation.email.hashicorp.com> to me 3:45 PM (0 minutes ago)

 HashiCorp

Please verify your email address

Press the button below to verify cloudtechintellugu@gmail.com for use with your HashiCorp account.

[Verify](#)

Or verify using this link:
<https://auth.hashicorp.com/u/email-verification?ticket=HIIxU9vNEHZQHzEEOlsVhz3vhBcSCDEh#>

If you feel this was an error, you can ignore this email.

Upgrade →



Service agreements

[Terms of Service](#) Required

I accept the Terms of Service

[Privacy Policy](#) Required

I accept the Privacy Policy

Marketing and newsletters

I would like to be updated via email about HashiCorp products, features, events, announcements, and education materials.

[Continue](#)

auth.hashicorp.com/u/signup/password?state=hKFo2SBueEZOM3lscnU3OG1SdWdjRnNHU1o4S25kRUZlWVIEZKFur3VuaXZlcNhbC1sb2dpbqN0aWTZlE1DZFESTUXMEFBcFIOZTFzbG42ZF4c09SU2ZPOXVHo2NpZNkgcEpLZ2JGdD)xZ...

Create HashiCorp Account

Enter your password to continue

cloudtechintelugu@gmail.com [Edit](#)

Password* [Reset](#)

.....

Your password must contain:

- ✓ At least 8 characters
- ✓ At least 3 of the following:
 - ✓ Lower case letters (a-z)
 - ✓ Upper case letters (A-Z)
 - ✓ Numbers (0-9)
 - ✓ Special characters (e.g. !@#\$%^&*)

[Create account](#)

Already have an account? [Log in](#)



Create organization

A HashiCorp Cloud Platform (HCP) organization allows you to deploy and manage HCP services and invite collaborators.

How will you use this organization? Not editable after creation



Business

- Best for company or institution use (even if this is just for testing for now).
- Ensures proper handling of taxes.
- Ensures business continuity by tying this organization to the business rather than individuals.



Personal

- Best for individual use for your own personal projects that have no affiliation to a business.

HCP organization name

- Must be unique
- May only include ASCII letters, numbers, dashes (-), and underscores (_). No spaces or special characters.

19 characters remaining

Create organization

Dashboard | HashiCorp Cloud +

portalcloud.hashicorp.com/orgs/b361d4b4-d67b-4111-8eaa-cfd054fe4b70/projects/59525d48-67c6-4713-b7a2-da7d92d1f8de

HashiConf 2025
Tickets are on sale now for HashiConf, HashiCorp's global cloud conference. Join us in San Francisco 9/24-26. [Register today](#)

default-project

Services
Boundary
Consul
Packer
Terraform
Vagrant Registry
Vault Dedicated
Vault Radar Trial
Waypoint
default-project
Active resources
HashiCorp Virtual Networks
Access control (IAM)
Project settings

cloudtechintelugu-org / Projects / default-project

default-project

Project dashboard
An overview of all resources in the project. [Learn more](#)

Project status: Active | Last resource seen: — | View project settings

Active resources

No active resources
There are 0 active resources inside this project.
[View active resources](#)

Billing summary

To view billing information across projects, go to the organization.
[View organization billing summary](#)

Boundary

Get started with Boundary

Deploy Boundary to access any system from anywhere based on user identity. Secure your first host in 15 minutes.
[Get started with Boundary](#)

Consul

Get started with Consul

Deploy your first Consul cluster and start building your service mesh.
[Get started with Consul](#)

Packer

Get started with Packer

Deploy your first Packer registry to better manage builds across cloud providers.
[Get started with Packer](#)

Terraform

Get started with Terraform

Infrastructure automation to provision and manage resources in any cloud or data center.
[Get started with Terraform](#)



Terraform

HCP Terraform signup

 HCP Terraform

Sign in to HCP Terraform

 Continue with HCP account

OR

Username or email

Password

[Forgot password?](#)

 Sign in

[Sign in with Terraform SSO.](#)

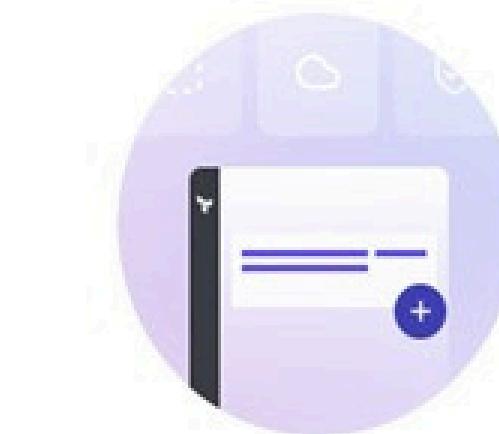
Need to sign up? Create your [free account](#).

View [Terraform Offerings](#) to find out which one is right for you.



Organizations

Terraform organizations let you manage organizations, projects, and teams.



Create your first organization

+ Create organization

New Organization | HCP Terraform

app.terraform.io/app/organizations/new

Organizations / New

Create a new organization

Organizations are privately shared spaces for teams to collaborate on infrastructure. [Learn more](#) about organizations in HCP Terraform.

How will you use this organization? Not editable after creation

 **Business**

- Best for company or institution use (even if this is just for testing for now).
- Ensures proper handling of taxes.
- Ensures business continuity by tying this organization to the business rather than individuals.

 **Personal**

- Best for individual use for your own personal projects that have no affiliation to a business.

Terraform organization name

- Must be unique.
- May contain valid characters including ASCII letters, numbers, spaces, as well as dashes (-), and underscores (_).

Email address

The organization email is used for any future notifications, such as billing alerts, and the organization avatar, via [gravatar.com](#).

Create organization **Cancel**

Support Terms Privacy Security Accessibility  © 2025 HashiCorp



Cost Estimation policy:

- we can estimate the cost of infra before we create it.
- By default, this feature is disable, we need to enable.
- Click on terraform logo -- > Organization -- > settings -- > Cost Estimation -- > enable

Sentinel Policy

- By this policy we can write our own conditions.
- It will check the condition of resource before it created.
- Resource will be created/ updated only when the condition is satisfied.
- EX: TAGS, VERIFIED AMIS, SG etc...



Terraform cloud features:

1. Workspaces
2. Projects
3. Runs
4. Variables and Variable Sets
5. Policies and Policy Sets
6. Run Tasks
7. Single Sign-On (SSO)
8. Remote State

TERRAFORM CLOUD OFFERS:

Free

UP TO
500 resources
per month

Cloud

Get started with all capabilities needed for infrastructure as code provisioning.

No credit card required

[Start for free](#)

Standard

STARTING AT
\$0.00014
per hour per resource

Cloud

For professional individuals or teams adopting infrastructure as code provisioning.

Enterprise support included

[Start for free](#)

Plus

Custom

Cloud

For enterprises standardizing and managing infrastructure automation and lifecycle, with scalable runs.

Enterprise support included

[Contact sales](#)

Enterprise

Custom

Self-managed

For enterprises with special security, compliance, and additional operational requirements.

Enterprise support included

[Contact sales](#)

First 500 resources per month are free



Secure Secrets in the Terraform Code:

- **Tip1:** Never store secrets in plain text
- **Tip2:** Mark Variables as sensitive
- **Tip3:** Use Environment variables
- **Tip4:** Secret Store (e.g.: Vault, AWS Secrets manager)
- **Note:** Even after marked a value as sensitive in tf file, they are stored within the Terraform state file.

Thank you